

Mihir Deshpande

+91 9623179829 Pune Email Blog LinkedIn GitHub

Software Engineer with 4 years of experience designing and developing scalable, high-performance backend systems. Skilled in integrating LLMs into production applications, with practical experience in prompt engineering and GenAI implementation. Adept at leading projects, improving system reliability, and collaborating cross-functionally to deliver solutions.

EDUCATION

Pune Institute of Computer Technology

B.E. Electronics and Telecommunication — CGPA: 9.15/10

Pune, India

Aug '17 - May '21

- Data structures and Algorithms, Object Oriented Programming, Operating Systems, Computer Networks, Machine Learning

WORK EXPERIENCE

Software Engineer

HelpShift

Pune, India

Mar '23 - Present

- Architected and optimized a scalable architecture for integrating a third-party LLM into a chatbot service handling 1.5K RPS, to achieve a **99.99% uptime**, reduce latency by **30%**, and ensure fault tolerance.
- Designed and developed a Distributed Asynchronous Task Queue, leveraging Kafka and Redis that **reduced system latency by 80% and load average by 30%** through offloading I/O and compute
- Contributed to a database migration from MongoDB to YugabyteDB, implementing a **dual-write strategy to maintain zero downtime**, resulting in a **60% improvement** in query performance and seamless data consistency
- Mentored **10+** new engineers, providing guidance on best coding practices, code reviews, and troubleshooting techniques, leading to faster onboarding, improved code quality, and increased team productivity
- Resolved **20+** on-call incidents and collaborated with SRE teams to implement auto-scaling policies, reducing overnight CPU spikes by **35%**.

Software Engineer

miniOrange Inc.

Pune, India

Aug '21 - Feb '23

- Enhanced system security by enforcing OAuth 2.0, JWT-based authentication, and multi-factor authentication (MFA), reducing unauthorized access incidents.
- Optimized Hibernate queries and indexed critical authorization tables, reducing database query execution time by **30%**, improving role lookup speed.

SKILLS

Programming Languages: Java, Clojure, Python

Databases and Streaming: MongoDB, PostgreSQL, Apache Kafka, Redis, Elasticsearch, YugabyteDB

Frameworks and Libraries: Spring Boot, Django

Core Competencies : Product Development, LLMs, Technical Leadership, Functional Programming, Distributed Systems

PROJECTS

Clj-deepwalk

Clojure, Graphs, Machine Learning

[View Project](#)

- A Graph Embedding Library that converts complex Networks into Vector Embeddings using Cosine Similarity, for Machine Learning analysis. Recognized in [Clojure Deref](#)

MCP Client for Ollama

Python, Ollama, MCP

[View project](#)

- Extensible MCP client for Ollama models
- Support for Git operations through MCP Git server, with interactive command-line assistant interface

AIde

Python, OpenAI, OOP, Prompt Engineering

[View project](#)

- An Intelligent code practice tool that harnesses OpenAI to dynamically generate coding challenges, evaluate solutions, and provide detailed feedback across multiple difficulty levels.

PUBLICATIONS

3D Modelling using Generative Adversarial Networks

Python, Keras, Deep Learning, Tensorflow

[View Paper](#)

- Resolved limitations in existing 3D GANs and engineered a novel solution for generating varied 3D models of target objects, enhancing the efficiency of 3D modeling processes